

Advanced Set Representation Methods

AVL trees. 2-3(-4) Trees.
Union-Find Set ADT

Advanced Set Representation. AVL Trees

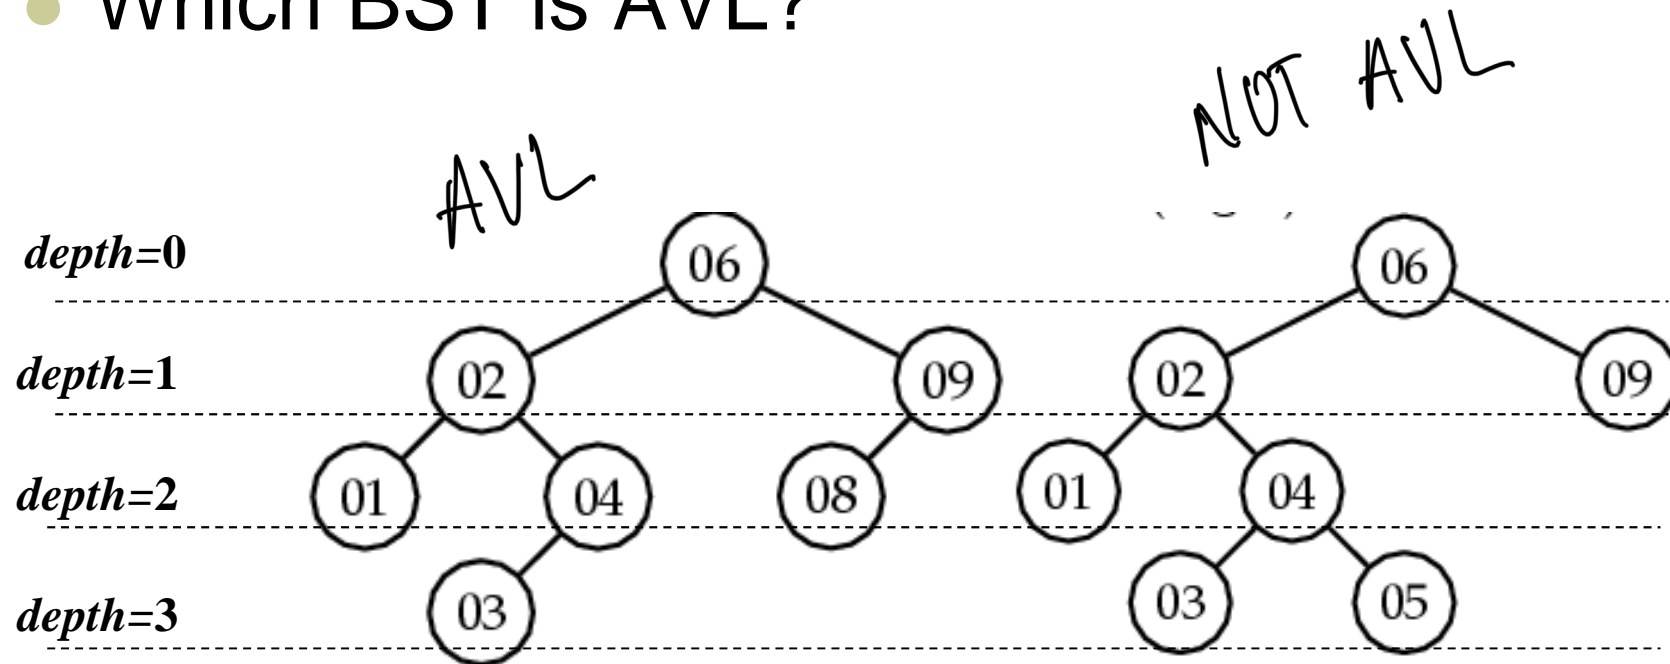
- Problem with BSTs: worst case operation may take $O(n)$ time. One solution:
- **AVL tree**: binary search tree with a *balance condition*:

For every node in an AVL tree T , the height of the left (T_L) and right (T_R) subtrees can differ by at most 1: $|h_L - h_R| \leq 1$

- Balance condition must be easy to maintain, and it ensures that the tree depth is $O(\log n)$
- Requiring that the left and right subtrees have the same height does not suffice (tree may not be shallow)

Two BSTs

- Which BST is AVL?




For every node in an AVL tree T , the height of the left (T_L) and right (T_R) subtrees can differ by at most 1: $|h_L - h_R| \leq 1$

AVL Tree Height

- Claim: height of an AVL tree storing n keys is $O(\log n)$
- Proof
 - Let $n(h)$ denote the number of nodes of an AVL tree of height h .
 - It is easy to see that $n(1) = 1$ and $n(2) = 2$
 - For $n > 2$, an AVL tree of height h contains:
 - the root node,
 - one AVL subtree of height $h-1$ and
 - another of height $h-2$.
 - That is, $n(h) = 1 + n(h-1) + n(h-2)$
 - Knowing $n(h-1) > n(h-2)$, we get $n(h) > 2n(h-2)$. By induction $n(h) > 2n(h-2)$, $n(h) > 4n(h-4)$, $n(h) > 8n(h-6)$, ..., $n(h) > 2^i n(h-2i)$
 - Solving the base case we get: $n(h) > 2^{h/2-1}$
 - Taking logarithms: $h < 2\log n(h) + 2$
- Thus the height of an AVL tree is $O(\log n)$

Insertion in an AVL Tree

- Insertion is as in a binary search tree
- Always done by attaching an external node.
- Example:




- After insertion the tree may need rebalancing

AVL tree rebalancing


- Two kinds of rotation:
 - single rotation (left or right)
 - double rotation
 - **left-right:** *left rotation* around the *left child* of a node followed by a *right rotation* around the *node itself*
 - **right-left:** *right rotation* around the *left child* of a node followed by a *left rotation* around the *node itself*
- Rotation features:
 - Nodes not in the subtree of the node rotated are unaffected
 - A rotation takes constant time
 - Before and after the rotation tree is still BST
 - Code for left rotation is symmetric to code for a right rotation

AVL Single Rotations


- Single Rotation
 - $k_1 < k_2$
 - all elements in subtree A are smaller than k_1
 - all elements in subtree C are larger than k_2
 - all elements in subtree B are in between k_1 and k_2




Single Rotation Right Example



- Note that subtrees B and C are empty



Single Rotation Left Example



AVL Implementation Detail

```
typedef struct
{
    ElementT element;
    AVLPtr left;
    AVLPtr right;
    int height;
} AVLNode;
typedef AVLNode *AVLPtr;
```

AVL Single Rotations

```
void snglRotRight( AVLPtr *k2 )
```

```
{
```

```
    AVLPtr k1 ;
```

```
    k1 = (*k2)->left ;
```

```
    (*k2)->left = k1->right ;
```

```
    k1->right = *k2 ;
```


```
    (*k2)->height = max( height(
        (*k2)->left ),
        height( (*k2)->right)
    ) + 1 ;
```

```
    k1->height = max( height( k1-
        >left ),
        (*k2)->height ) + 1 ;
```



```
    *k2 = k1 ; /* assign new root */
```

```
}
```

```
/* snglRotLeft is symmetric */
```




Single Rotation not enough



AVL Double Rotation


- Right-left
 - $k_2 < k_1, k_3 < k_1, k_3 < k_2$
 - right rotation around the left child of a node followed by
 - a left rotation around the node itself
 - Rotate to make k_2 topmost node

```
void dblRotLeft( AVLPtr k3 )  
{  
    /* rotate between k1 and k2 */  
    snglRotRight ( (*k3)->left );  
    /* rotate between k3 and k2 */  
    snglRotLeft ( k3 );  
}
```



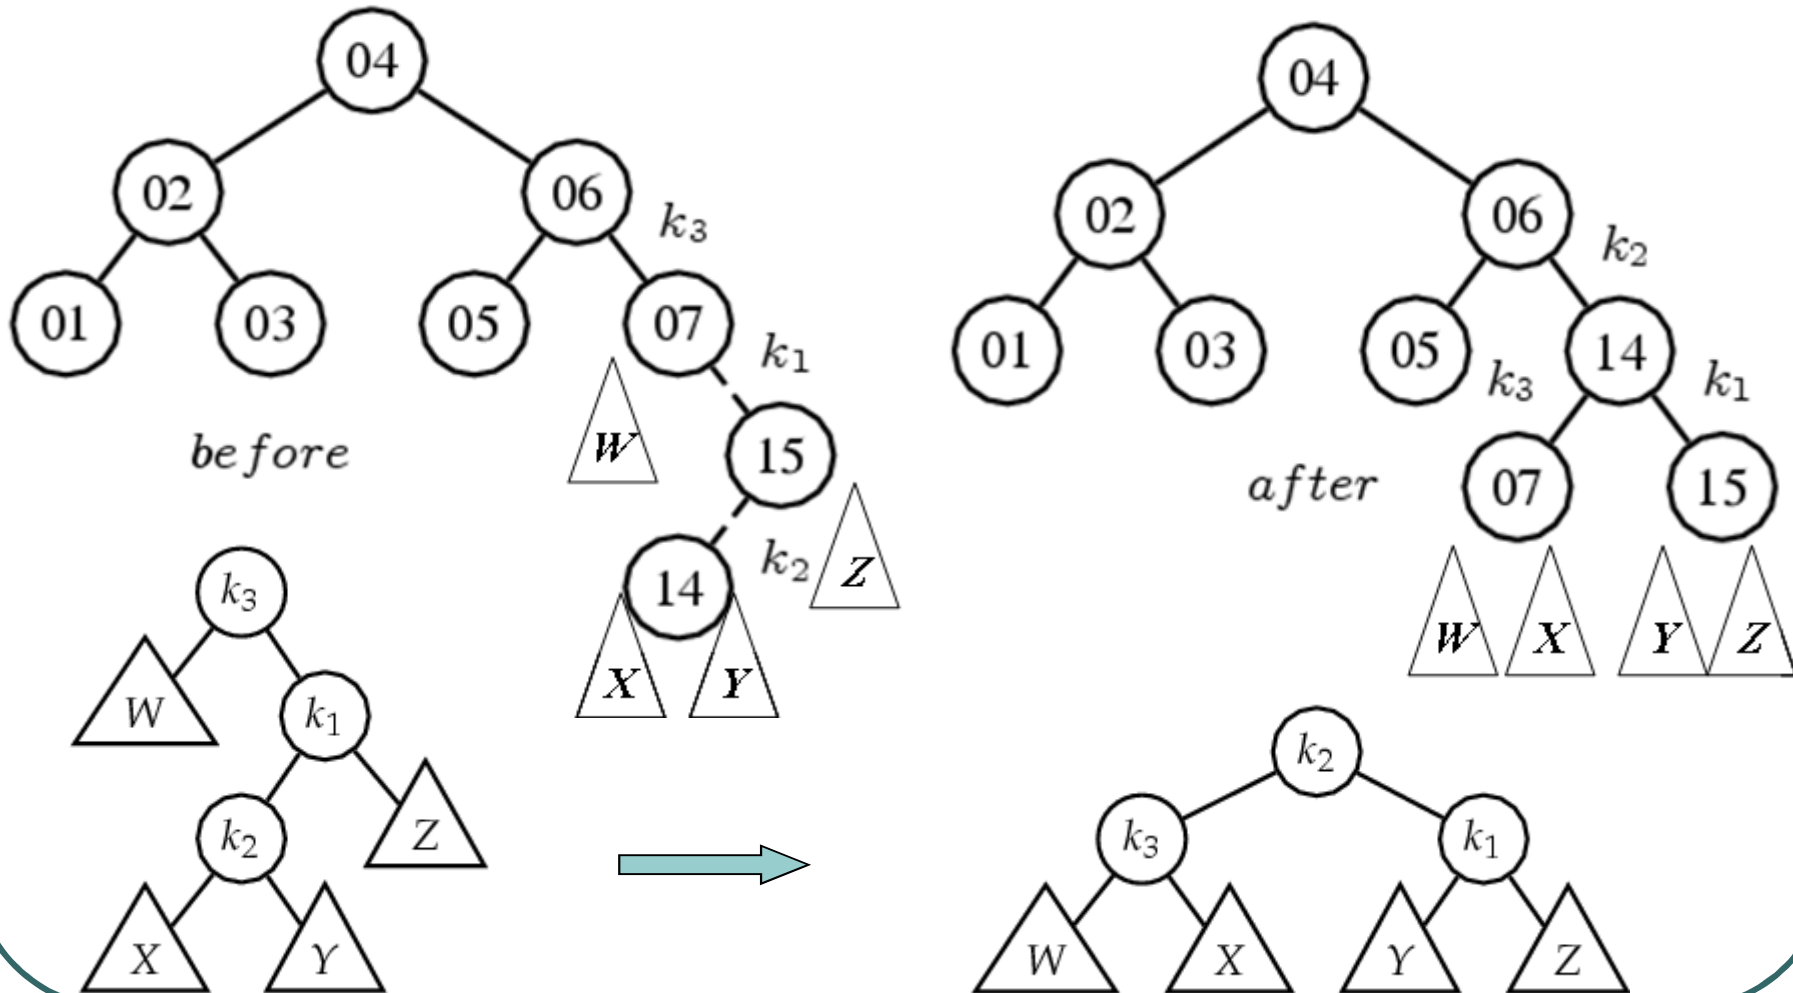
AVL Double Rotation

- Left-right:
 - $k_1 < k_2, k_1 < k_3, k_2 < k_3$
 - left rotation around the left child of a node followed by a
 - right rotation around the node itself
 - Rotate to make k_2 topmost node




AVL Right-left Double Rotation Example 1

- Subtrees W, X, Y, Z are empty




AVL Right-left Double Rotation Example 2

- Subtree X is empty




AVL Double Rotation example

- A symmetric case of double rotation



Tree before
insertion of 27

Tree after
insertion of 27
and re-balancing



AVL Trees. When to use rotations

- Balance factor: $height(T_L) - height(T_R)$
- Use
 - **Single rotation left:** when a node is inserted in the right subtree of the right child (B) of the nearest ancestor (A) with balance factor -2
 - **Single rotation right:** when a node is inserted in the left subtree of the left child (B) of the nearest ancestor (A) with balance factor +2.
 - **Right-left double rotation:** when a node is inserted in the left subtree of the right child (B) of the nearest ancestor (A) with balance factor -2.
 - **Left-right double rotation:** when a node is inserted in the right subtree of the left child (B) of the nearest ancestor (A) with balance factor +2.

AVL Trees. Deletions

- A node is deleted using the standard inorder successor (predecessor) logic for binary search trees
- Imbalance is fixed using rotations
- Identify the parent of the actual node that was deleted, then:
 - If the left child was deleted, the balance factor at the parent decreased by 1.
 - If the right child was deleted, the balance factor at the parent increased by 1.

AVL Trees. Deletions. Rebalancing (I)

- Let A be the node where balance must be restored.
- If the deleted node was in A's right subtree, then let B be the root of A's left subtree. Then:
 - **B has balance factor 0 or +1 after deletion** -- then perform a single right rotation
 - **B has balance factor -1 after deletion** -- then perform a left-right rotation
- If the deleted node was in A's left subtree, then let B be the root of A's right subtree. Then:
 - **B has balance factor 0 or -1 after deletion** -- then perform a single left rotation
 - **B has balance factor +1 after deletion** -- then perform a right-left rotation

AVL Trees. Deletions. Rebalancing (II)

- Unlike insertion, *one rotation may not be enough to restore balance to the tree*. If this is the case, then locate the next node where the balance factor is "bad" (call this A):
- If A's balance factor is positive, then let B be A's left child
 - If B's left subtree height is larger than B's right subtree height, then perform a **single right rotation**.
 - If B's left subtree height is smaller than B's right subtree height, then perform a **left-right rotation**.
 - If B's left subtree height is equal to B's right subtree height, then perform **either rotation**.

AVL Trees. Deletions. Rebalancing (III)

- If A's balance factor is negative, then let B be A's right child
 - If B's right subtree height is larger than B's left subtree height, then perform a **single left rotation**.
 - If B's right subtree height is smaller than B's left subtree height, then perform a **right-left rotation**.
 - If B's right subtree height is equal to B's left subtree height, then perform **either rotation**.

AVL trees demo

- Demos from:

<http://webpages.ull.es/users/jriera/Docencia/AVL/AVL%20tree%20applet.htm>

<http://www.site.uottawa.ca/~stan/csi2514/applets/avl/BT.html>

- Replacement for AVL trees: Red-Black trees (not discussed here)

Running times for AVL tree operations


- a single restructure is $O(1)$
 - using a linked-structure binary tree
- find is $O(\log n)$
 - height of tree is $O(\log n)$, no restructures needed
- insert is $O(\log n)$
 - initial find is $O(\log n)$
 - restructuring up the tree, maintaining heights is $O(\log n)$
- remove is $O(\log n)$
 - initial find is $O(\log n)$
 - restructuring up the tree, maintaining heights is $O(\log n)$

2-3 Trees

- 2-3 tree properties:
 - Each interior node has two or three children.
 - Each path from the root to a leaf has the same length.
 - A tree with zero or one node(s) is a special case of a 2-3 tree.
- Representing sets with 2-3 trees:
 - Elements are placed at the leaves
 - If element a is to the left of element b , then $a < b$ must hold.
 - Ordering of elements based on one field of a record: a **key**.
 - At each **interior** node: key of the smallest descendant of the second child and, if there is a third child, key of the smallest descendant of third child.

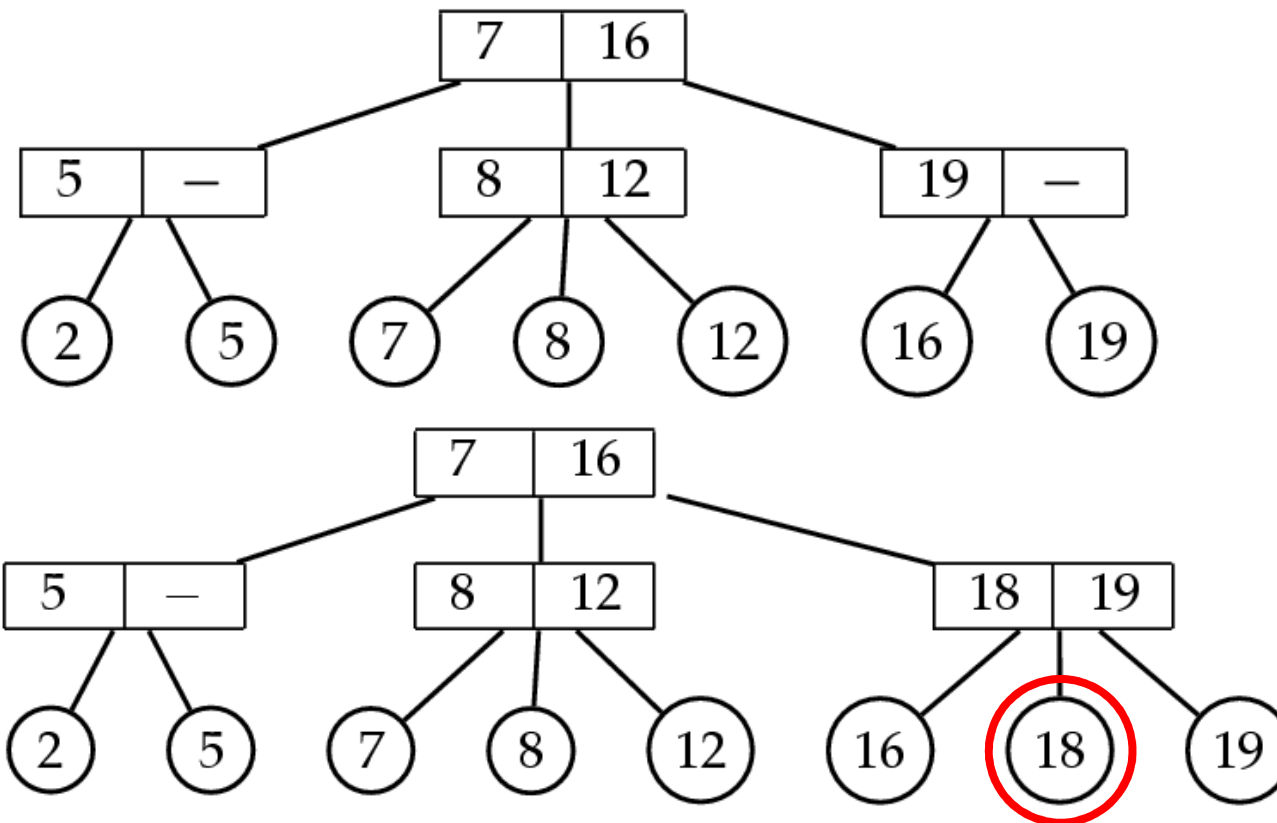
2-3 trees

- A 2-3 tree of k levels has between 2^{k-1} and 3^{k-1} leaves, i.e.
 - representing a set of n elements requires at least $1 + \log_3 n$ levels and no more than $1 + \log_2 n$ levels. Thus, path lengths in the tree are $O(\log n)$.



2-3 trees. Insertion – 2 children

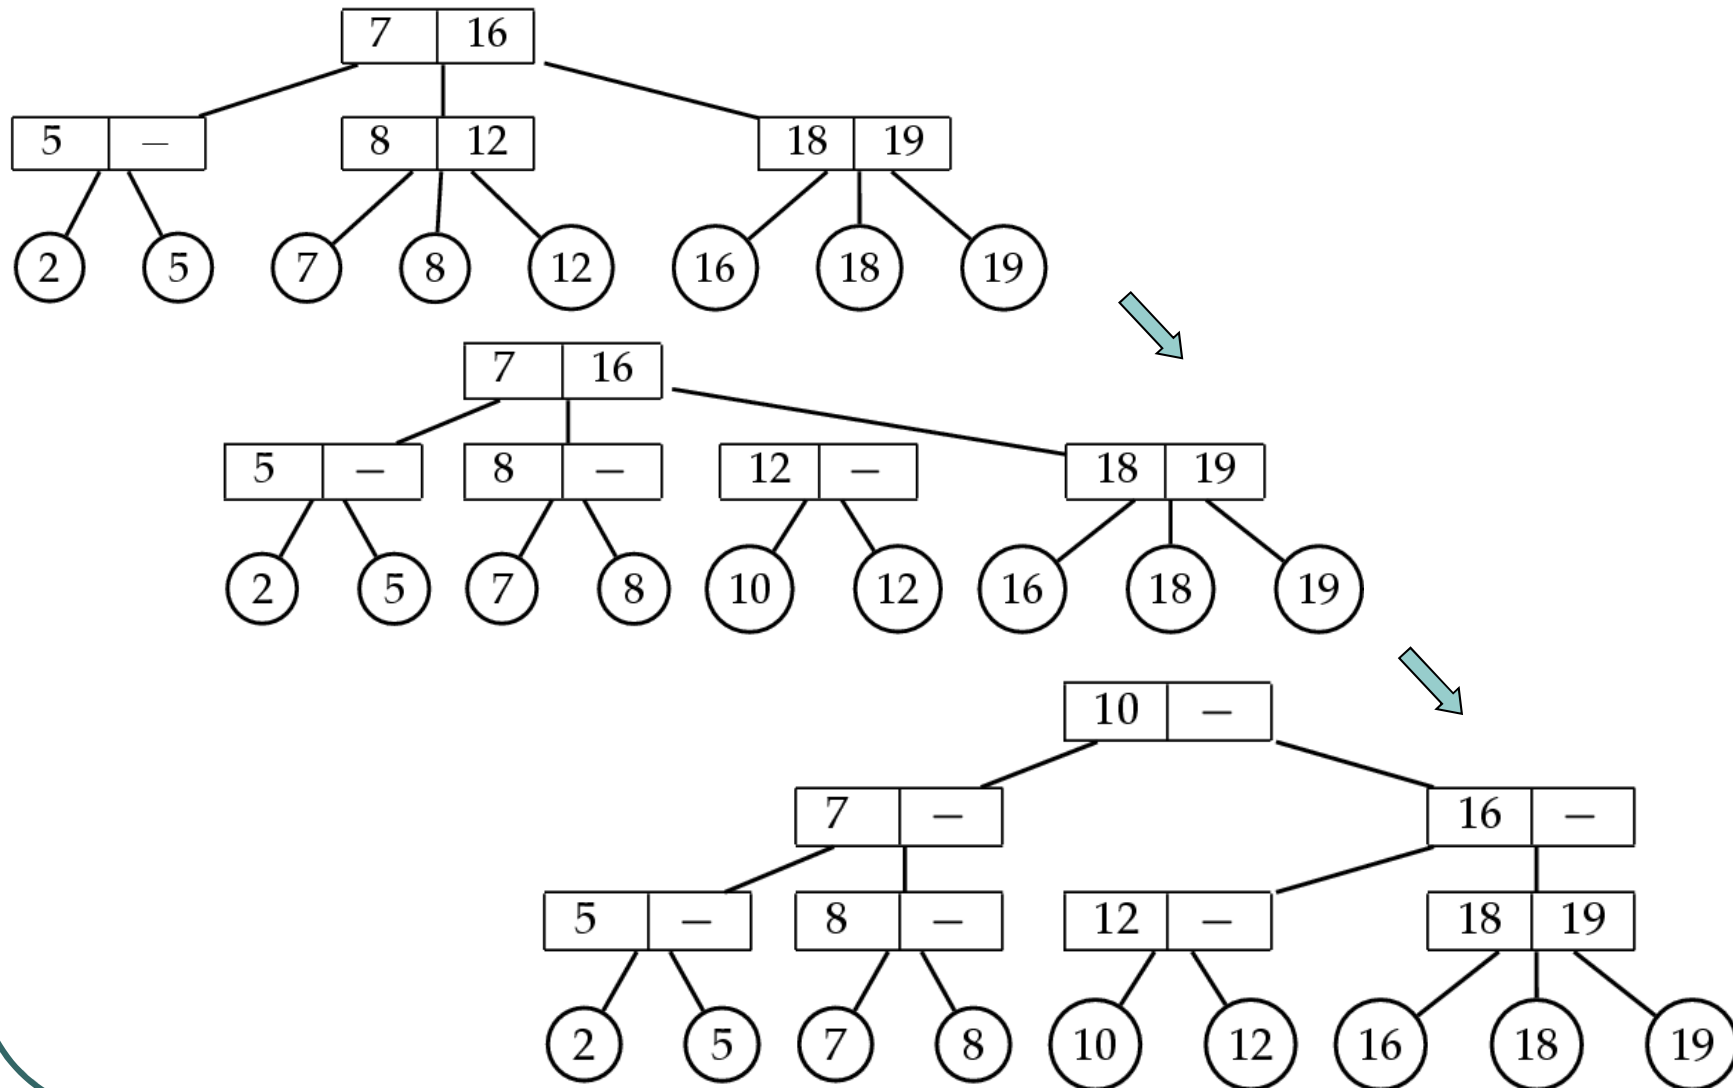
- If parent node has only two children:
 - insert in order, adjust keys



2-3 trees. Insertion – 3 children

- If parent node *node* already has three children:
 - We cannot have four nodes, and thus we have to split parent in two nodes *node* and *node'*.
 - The two smallest elements among the four children of *node* stay with *node*,
 - The two larger will become the children of *node'*
 - Continue process up the tree
 - Special case when splitting the root

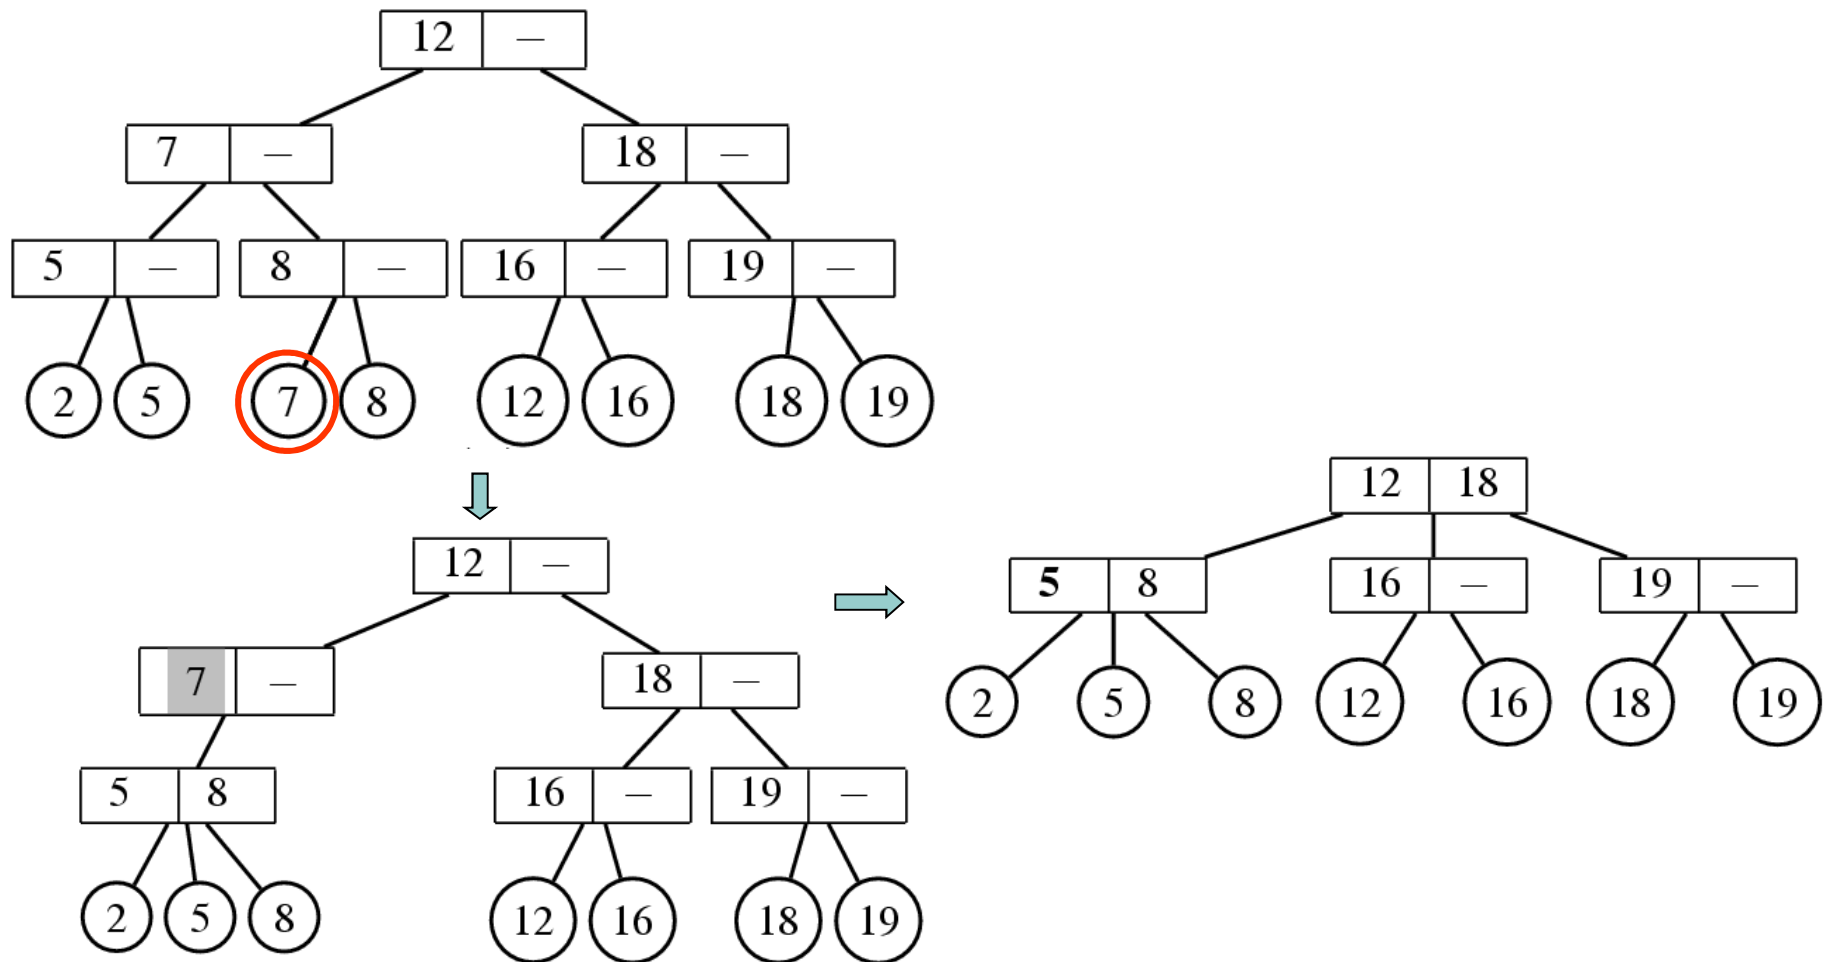
2-3 trees. Insertion



2-3 trees. Deletion

- Leaf deletion could leave parent with only 1 child
 - Parent=root: delete *node* and let its lone child be the new root
 - Otherwise, let p be the parent of *node*
 - if p has another child, and that child of p has three children, then we can transfer the proper one to *node*; done;
 - if the children of p , have only two children, transfer the lone child of *node* to an adjacent sibling of *node*, and delete *node*;
 - p has only one child, repeat all the above, recursively, with p in place of *node*.




2-3 trees. Deletion



2-3-4 trees.


- 2-3-4 tree refer to how many links to child nodes can potentially be contained in a given node. For non-leaf nodes, three arrangements :
 - A node with one data item always has two children
 - A node with two data items always has three children
 - A node with three data items always has four children
 - In short, a non-leaf node must always have one more child than it has data items.
- Symbolically, if the number of child links is L and the number of data items is D , then $L = D + 1$
- Empty nodes are not allowed.

2-3-4 trees. Node kinds.




2-3-4 trees. Properties and example.

- All children in the subtree rooted at child 0 have key values less than key 0.
- All children in the subtree rooted at child 1 have key values greater than key 0 but less than key 1.
- All children in the subtree rooted at child 2 have key values greater than key 1 but less than key 2.
- All children in the subtree rooted at child 3 have key values greater than key 2.



2-3-4 trees. Example.



2-3-4 Trees. Insertion

- **Insertion procedure:**

- similar to insertion in 2-3 trees
- items are inserted at the leafs
- since a 4-node cannot take another item, 4-nodes are split up during insertion process

- **Strategy**

- on the way from the root down to the leaf:
split up all 4-nodes "on the way"
- → insertion can be done in one pass
(remember: in 2-3 trees, a reverse pass might be necessary)

2-3-4 Trees. Deletion

- **Deletion procedure:**
 - similar to deletion in 2-3 trees
 - items are deleted at the leafs
 - swap item of internal node with inorder successor
 - note: a 2-node leaf creates a problem
- **Strategy** (different strategies possible)
 - on the way from the root down to the leaf:
turn 2-nodes (except root) into 3-nodes
 - → deletion can be done in one pass
(remember: in 2-3 trees, a reverse pass might be necessary)

Demo: 2-3-4 trees

- Demo from:
<http://www.cs.unm.edu/~rlpm/499/ttft.html>
- Local:
[Demos\2-3-4trees\TTFT.jar](#)

Disjoint Sets with the UNION and FIND Operations

- Applicable to problems where:
 - start with a collection of objects, each in a set by itself;
 - combine sets in some order, and from time to time
 - ask which set a particular object is in
- Equivalence classes:
 - If set S has an equivalence relation (reflexive, symmetrical, transitive) defined on it, then the set S can be partitioned into disjoint subsets S_1, S_2, \dots, S_k with $\bigcup_k S_k = S$
- Equivalence problem:
 - given a set S and a sequence of statements of the form $a \equiv b$
 - process the statements in order in such a way that at any time we are able to determine in which equivalence class a given element belongs

Union-find set ADT

- Example statements (see previous slide)

$$S = \{1, 2, \dots, 7\}$$

$$1 \equiv 2, 5 \equiv 6, 3 \equiv 4, 1 \equiv 4$$

$$1 \equiv 2 \quad \{1, 2\} \{3\} \{4\} \{5\} \{6\} \{7\}$$

$$5 \equiv 6 \quad \{1, 2\} \{3\} \{4\} \{5, 6\} \{7\}$$

$$3 \equiv 4 \quad \{1, 2\} \{3, 4\} \{5, 6\} \{7\}$$


$$1 \equiv 4 \quad \{1, 2, 3, 4\} \{5, 6\} \{7\}$$

- Operations:
 - union(A, B) takes the union of the components A and B and calls the result either A or B, arbitrarily.
 - find(x), a function that returns the name for the component of which x is a member.
 - initial(A, x) creates a component named A that contains only the element x.


Union-find set ADT. Implementations

- List-based
 - Each set is stored in a sequence represented with a linked-list
 - Each node should store an object containing the element and a reference to the set name
- Tree based
 - Each element is stored in a node, which contains a pointer to a set name
 - A node v whose set pointer points back to v is also a set name
 - Each set is a tree, rooted at a node with a self-referencing set pointer


Examples



List based




Tree based



Union-find set. List based implementation

```
const int n = /* appropriate value */
typedef int NameT;
typedef int ElementT;
typedef struct ufset
{
    struct /* headers for set lists */
    {
        int count;;
        int firstElement ;
    } setHeaders[ n ] ;
```

```
    struct /* table giving set containing
        each member */
    {
        NameT setName ;
        int nextElement ;
    } setNames[ n ] ;
} UFSetT;
```



Union-find sets. Tree representation

- Maintain S as a collection of trees (a forest), one per partition
- Initially, there are n trees, each containing a single element
- $\text{find}(i)$ returns the root of the tree containing i
- $\text{union}(i, j)$ merges the trees containing i and j
- Typical traversals not needed, so no need for pointers to children, but we need pointers to parent
- Parent pointers can be stored in an array: $\text{parent}[i]$ ($=-1$ for root)

Union-find sets. Tree implementation

INIT()

```
1  for  $i \leftarrow 1$  to  $n$ 
2      do  $parent[i] \leftarrow 0$ 
```

FIND(i)


```
1   $j \leftarrow i$ 
2  while  $parent[j] > 0$ 
3      do  $j \leftarrow parent[j]$ 
4  return  $j$ 
```

UNION(i, j)

```
1   $root_1 \leftarrow FIND(i)$ 
2   $root_2 \leftarrow FIND(j)$ 
3  if  $root_1 \neq root_2$ 
4      then  $parent[root_2] \leftarrow root_1$ 
```


Union-find sets. Speeding up

- Union by size (rank):
 - When performing a union, make the root of smaller tree point to the root of the larger
- Implies $O(n \log n)$ time for performing n unionfind operations:
 - Each time we follow a pointer, we are going to a subtree of size at least double the size of the previous subtree
 - Thus, we will follow at most $O(\log n)$ pointers for any find.



Union-find sets. Speeding up

- Path compression:
 - After performing a find, compress all the pointers on the path just traversed so that they all point to the root
- Implies $O(n \log^* n)$ time for performing n union- find



Using union by rank and path compression

MAKESET(x)

```
1   $p[x] \leftarrow x$   
2   $r[x] \leftarrow 0$ 
```

FINDSET(x)

```
1  if  $x \neq p[x]$   
2      then  $p[x] \leftarrow \text{FINDSET}(p[x])$   
3  return  $p[x]$ 
```

LINK(x, y)

```
1  if  $r[x] > r[y]$   
2      then  $p[y] \leftarrow x$   
3      else  $p[x] \leftarrow y$   
4          if  $r[x] = r[y]$   
5              then  $r[y] \leftarrow r[y] + 1$ 
```

UNION(x, y)

```
1  LINK(FINDSET( $x$ ), FINDSET( $y$ ))
```

- Time bound of $O(m\alpha(n))$ where $\alpha(n)$ is a very slowly growing function

Union-find ADT demo

- Demo from:
<http://www.cs.unm.edu/~rlpm/499/uf.html>
- Local:
[Demos\UnionFind\UnionFind.jar](#)

A very quickly growing function and its inverse

- For integers $k \geq 0$ and $j \geq 1$, define $A_k(j)$:

$$A_k(j) = \begin{cases} j+1, & \text{if } k = 0 \\ A_{k-1}^{(j+1)}(j), & \text{if } k \geq 1 \end{cases}$$

- where $A_{k-1}^0(j) = j$, $A_{k-1}^{(i)}(j) = A_{k-1}(A_{k-1}^{(i-1)}(j))$ for $i \geq 1$.
- k is called the level of the function and
- i in the above is called iterations.
- $A_k(j)$ strictly increase with both j and k .
- Let us see how quick the increase is

A very quickly growing function and its inverse

- Let us see $A_k(1)$: for $k=0,1,2,3,4$.
 - $A_0(1)=1+1=2$
 - $A_1(1)=2\times 1+1=3$
 - $A_2(1)=2^{1+1}\times(1+1)-1=7$
 - $A_3(1)=A_2^{(1+1)}(1)=A_2^{(2)}(1)=A_2(A_2(1))=A_2(7)=2^{7+1}(7+1)-1=2^8\times 8-1=2047$
 - $A_4(1)=A_3^2(1)=A_3(A_3(1))=A_3(2047)=A_2^{(2048)}(2047)$
 - $\gg A_2(2047)=2^{2048}\times 2048-1 > 2^{2048}=(2^4)^{512}=(16)^{512}$
 - $\gg \mathbf{10^{80}}$. (estimated number of atoms in universe)

Inverse of $A_k(n)$: $\alpha(n)$

- $\alpha(n) = \min\{k: A_k(1) \geq n\}$
- $\alpha(n) = 0$ for $0 \leq n \leq 2$
1 for $n = 3$
2 for $4 \leq n \leq 7$
3 for $8 \leq n \leq 2047$
4 for $2048 \leq n \leq A_4(1)$.
- Extremely slow increasing function.
- $\alpha(n) \leq 4$ for all practical purpose.

Reading

- AHU, chapter 5, sections 5.4, 5.5
- Preiss, chapter: Search Trees – section AVL Search Trees
- CLR, chapter 22, sections 1-4
- CLRS chapter 21
- Notes